

Loss function + SGD

Razvan Marinescu

What is a loss?

- Recall least square

$$J(\theta) = \frac{1}{2} \sum_{i=1}^3 (\mathbf{x}^{(i)} \cdot \theta - y^{(i)})^2 = \frac{1}{2} \sum_{i=1}^3 (\mathbf{x}^{(i)\top} \theta - y^{(i)})^2$$

- Logistic regression

$$\log L(\boldsymbol{\theta}) = \sum_{i=1}^m y^{(i)} \log(g(\boldsymbol{\theta}^\top \mathbf{x}^{(i)})) + (1 - y^{(i)}) \log(1 - g(\boldsymbol{\theta}^\top \mathbf{x}^{(i)}))$$

Linear regression – review

- Another way of writing this

$$L(\theta) := \sum_{i \in \text{data}} \ell(\theta; \mathbf{x}^{(i)}, y^{(i)})$$

- Empirical risk minimization
- How do we optimize it?

Loss functions

a = model prediction

y = label from training/test data

$$0/1: l^{(0/1)}(y, a) = 1[y a \leq 0]$$


$$\text{Hinge: } l^{(hin)}(y, a) = \max\{0, 1 - ya\}$$

$$\text{Logistic: } l^{(log)}(y, a) = \frac{1}{\log 2} \log(1 + \exp[-ya])$$

$$\text{Exponential: } l^{(exp)}(y, a) = \exp[-ya]$$


Different loss function can
result in different classifiers

These are all convex and so can be
minimized using algorithms like
Gradient Descent



Refresher: Gradient Descent

- Used for finding the minima of an optimization function
- Idea: Start at a random initial guess on the function and iteratively **move towards its minima**



How do we know in which direction is the minima?

Direction of negative gradient

***Refresher: gradient of a function points in the direction of the greatest rate of increase of the function, and its magnitude is the slope of the graph in that direction.

Stochastic gradient descent

a single data-point pair

- Randomly order data $\mathbf{x}^{(i)}, y^{(i)}$;
- Compute $\partial_{\theta} \ell(\theta_i; \mathbf{x}^{(i)}, y^{(i)})$;
- Update $\theta_{i+1} := \theta_i - \eta \cdot \partial_{\theta} \ell(\theta_i; \mathbf{x}^{(i)}, y^{(i)})$.



Repeat until
convergence

- Consider SGD for example $\mathbf{x}^{(1)} = (1, 2)$, $y^{(1)} = 5$, and θ initially $(1, 1)$.
- The squared-error on this example is $(1 * \theta_0 + 2 * \theta_1 - 5)^2 = 4$, and its contribution to $J(\theta)$ is 2 (half the squared error).

At $\theta = (1, 1)$:


$$\frac{\partial \text{contribution}}{\partial \theta_0} = \frac{2}{2}(1 * \theta_0 + 2 * \theta_1 - 5) * 1 = -2$$

$$\frac{\partial \text{contribution}}{\partial \theta_1} = \frac{2}{2}(1 * \theta_0 + 2 * \theta_1 - 5) * 2 = -4$$

- with step size $\frac{1}{20}$, update θ to $\theta - \frac{1}{20} \nabla_{\theta}(\text{contribution})$.
- New $\theta = (1, 1) - (\frac{-2}{20}, \frac{-4}{20}) = (1.1, 1.2)$

One ML model: Linear classification

- Outputs a **classification** (one of N possible classes) based on the value of a **linear combination of the characteristics** (features)




Example with 2 features – 2D feature vector (x_1, x_2) :

- Linear classifiers H_1 and H_2 successfully partition the two classes of dots
- H_3 does not

Which is better, H_1 or H_2 ? Why?

Goal: Find a linear decision boundary

Linear classification



$$\mathbf{x} \cdot \mathbf{w} = |\mathbf{x}| \times |\mathbf{w}| \times \cos \theta$$

$$\mathbf{x} \cdot \mathbf{w} = \mathbf{x}^T \mathbf{w} = \mathbf{w}^T \mathbf{x}$$

$$\mathbf{x}_0 \cdot \mathbf{w} = \mathbf{x}_1 \cdot \mathbf{w} = \mathbf{x}_2 \cdot \mathbf{w} = b$$

How to determine if a feature vector \mathbf{x} is on the $+$ or $-$ side of the line?

Evaluate the dot product of \mathbf{x} and \mathbf{w} :


- If $\mathbf{x} \cdot \mathbf{w} > b$, then $+$
- Otherwise $-$

2 features means **2D classification** and a **1D classification boundary**

N features means **N dimensional classification** and an **N-1 dimensional classification boundary**

Dimensions	Linear boundary
1	Point
2	Line
3	Plane
>3	Hyperplane

Classifier geometry – w and b



Non-homogeneous:

$$\mathbf{w}^T \mathbf{x} - b = [w_1 \quad w_2] \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} - b > 0$$

Homogeneous:

$$\mathbf{w}^T \mathbf{x} = [w_1 \quad w_2 \quad -b] \begin{bmatrix} x_1 \\ x_2 \\ 1 \end{bmatrix} > 0$$

Is w a unit vector?

Doesn't have to be

What's the relationship
between w and b ?

$$(w, b) \equiv (kw, kb)$$


$$[w_1 \quad w_2] \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} - b = 0 \quad [2w_1 \quad 2w_2] \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} - 2b = 0$$

These describe the same line

Classifier geometry – w and b

Non-homogeneous:

$$w^T x - b = [w_1 \quad w_2] \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} - b > 0$$



Note:

Regular classifier: the **margin** is the distance from the decision boundary

Scoring classifier: the **margin** is the score


In both cases, value is **positive** for correctly classified, **negative** for incorrectly classified

Minimum Error Hyperplane

- Error of a linear model (\mathbf{w}, b) for an instance (\mathbf{x}_n, y_n)

$$\mathbf{1}[y_n(\mathbf{w} \cdot \mathbf{x}_n + b) \leq 0]$$

Indicator function: 1 if stuff inside is true (incorrect prediction) and 0 otherwise (correct prediction)




- Objective function to minimize to find the minimum error hyperplane:

$$\min_{\mathbf{w}, b} \sum_n \mathbf{1}[y_n(\mathbf{w} \cdot \mathbf{x}_n + b) \leq 0]$$

Minimum Error Hyperplane

- Error of a linear model (\mathbf{w}, b) for an instance (\mathbf{x}_n, y_n)

$$\mathbf{1}[y_n(\mathbf{w} \cdot \mathbf{x}_n + b) \leq 0]$$


 **Indicator function:** 1 if stuff inside is true (incorrect prediction) and 0 otherwise (correct prediction)

- Objective function to minimize to find the minimum error hyperplane:

$$\min_{\mathbf{w}, b} \sum_n \mathbf{1}[y_n(\mathbf{w} \cdot \mathbf{x}_n + b) \leq 0]$$

This optimization formulation is minimizing the **0/1 loss**


0/1 Loss



Minimizing 0-1 loss is difficult!

Why?


0/1 Loss



Minimizing 0-1 loss is difficult!

Why?

Not smooth, and not convex



Alternatives to 0/1 Loss

- Need to find an **upper-bound** to 0/1 loss that is **convex**
- Why do we require (1) convexity, and (2) upper-boundedness?
- (1) So that minimization is easy (we know how to minimize a convex function)
- (2) So that minimizing the upper bound also *pushes down* the real objective

Convex upper-bounds of 0/1 Loss

0/1: $l^{(0/1)}(y, a) = 1[y a \leq 0]$


Hinge: $l^{(hin)}(y, a) = \max\{0, 1 - ya\}$

Logistic: $l^{(log)}(y, a) = \frac{1}{\log 2} \log(1 + \exp[-ya])$

Exponential: $l^{(exp)}(y, a) = \exp[-ya]$

Different loss function can
result in different classifiers

These are all convex and so can be
minimized using algorithms like
Gradient Descent



Why does it work?

- Law of large number

$$\frac{1}{|\text{data}|} \sum_{i \in \text{data}} \ell(\theta; \mathbf{x}^{(i)}, y^{(i)}) \Rightarrow \mathbb{E}_{\text{distribution}}[\ell(\theta; \mathbf{x}, y)]$$

- Empirical risk minimization approximates the true optimal model
- Is the loss ℓ proper?

Classification calibration

- We say a loss function ℓ is Φ -calibrated if there exists a non-decreasing function Φ such that for all f

$$\Phi \left(R_\ell(f) - \min_{f'} R_\ell(f') \right) \leq R(f) - \min_{f'} R(f')$$

where in above

$$R_\ell(f) := \mathbb{E}[\ell(f; X, Y)], \quad R(f) := \mathbb{E}[1(f(X) \neq Y)]$$